

Coffee

Sean Cerulean Johnson

2022-04-03

Data : Importing and Cleaning

From TidyTuesday URL:<https://github.com/rfordatascience/tidytuesday/tree/master/data/2020/2020-07-07>

Note: within the above link, there was already some pre-processing done to the data with the column and value names.

Quick Overview Summary

```
summary(coffee_ratings)
```

```
## total_cup_points species owner country_of_origin
## Min. : 0.00 Length:1339 Length:1339 Length:1339
## 1st Qu.:81.08 Class :character Class :character Class :character
## Median :82.50 Mode :character Mode :character Mode :character
## Mean :82.09
## 3rd Qu.:83.67
## Max. :90.58
##
## farm_name lot_number mill ico_number
## Length:1339 Length:1339 Length:1339 Length:1339
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
##
##
##
## company altitude region producer
## Length:1339 Length:1339 Length:1339 Length:1339
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
##
##
##
## number_of_bags bag_weight in_country_partner harvest_year
## Min. : 0.0 Length:1339 Length:1339 Length:1339
```

```

## 1st Qu.: 14.0    Class :character    Class :character    Class :character
## Median : 175.0    Mode :character    Mode :character    Mode :character
## Mean : 154.2
## 3rd Qu.: 275.0
## Max. :1062.0
##
## grading_date      owner_1          variety          processing_method
## Length:1339      Length:1339      Length:1339      Length:1339
## Class :character  Class :character  Class :character  Class :character
## Mode :character  Mode :character  Mode :character  Mode :character
##
##
##
##
##      aroma      flavor      aftertaste      acidity      body
## Min. :0.000    Min. :0.00    Min. :0.000    Min. :0.000    Min. :0.000
## 1st Qu.:7.420    1st Qu.:7.33    1st Qu.:7.250    1st Qu.:7.330    1st Qu.:7.330
## Median :7.580    Median :7.58    Median :7.420    Median :7.580    Median :7.500
## Mean :7.567    Mean :7.52    Mean :7.401    Mean :7.536    Mean :7.517
## 3rd Qu.:7.750    3rd Qu.:7.75    3rd Qu.:7.580    3rd Qu.:7.750    3rd Qu.:7.670
## Max. :8.750    Max. :8.83    Max. :8.670    Max. :8.750    Max. :8.580
##
##      balance      uniformity      clean_cup      sweetness
## Min. :0.000    Min. : 0.000    Min. : 0.000    Min. : 0.000
## 1st Qu.:7.330    1st Qu.:10.000    1st Qu.:10.000    1st Qu.:10.000
## Median :7.500    Median :10.000    Median :10.000    Median :10.000
## Mean :7.518    Mean : 9.835    Mean : 9.835    Mean : 9.857
## 3rd Qu.:7.750    3rd Qu.:10.000    3rd Qu.:10.000    3rd Qu.:10.000
## Max. :8.750    Max. :10.000    Max. :10.000    Max. :10.000
##
## cupper_points      moisture      category_one_defects      quakers
## Min. : 0.000    Min. :0.00000    Min. : 0.0000    Min. : 0.0000
## 1st Qu.: 7.250    1st Qu.:0.09000    1st Qu.: 0.0000    1st Qu.: 0.0000
## Median : 7.500    Median :0.11000    Median : 0.0000    Median : 0.0000
## Mean : 7.503    Mean :0.08838    Mean : 0.4795    Mean : 0.1734
## 3rd Qu.: 7.750    3rd Qu.:0.12000    3rd Qu.: 0.0000    3rd Qu.: 0.0000
## Max. :10.000    Max. :0.28000    Max. :63.0000    Max. :11.0000
##                                     NA's :1
##      color      category_two_defects      expiration      certification_body
## Length:1339    Min. : 0.000    Length:1339    Length:1339
## Class :character  1st Qu.: 0.000    Class :character  Class :character
## Mode :character  Median : 2.000    Mode :character  Mode :character
##                                     Mean : 3.556
##                                     3rd Qu.: 4.000
##                                     Max. :55.000
##
## certification_address certification_contact unit_of_measurement
## Length:1339      Length:1339      Length:1339
## Class :character  Class :character  Class :character
## Mode :character  Mode :character  Mode :character
##
##
##
##

```

```
## altitude_low_meters altitude_high_meters altitude_mean_meters
## Min. : 1 Min. : 1 Min. : 1
## 1st Qu.: 1100 1st Qu.: 1100 1st Qu.: 1100
## Median : 1311 Median : 1350 Median : 1311
## Mean : 1751 Mean : 1799 Mean : 1775
## 3rd Qu.: 1600 3rd Qu.: 1650 3rd Qu.: 1600
## Max. :190164 Max. :190164 Max. :190164
## NA's :230 NA's :230 NA's :230
```

Quite a few NA's.

Numerical Columns: 1 within quakers, and 230 in Altitude low/high/mean.

Next, need to check what is happening in the rest of the data set, the character type.

Count of NA's per column

```
#type of data, col = 2, type of function applied
apply(X=is.na(coffee_ratings), MARGIN = 2, FUN = sum)
```

```
## total_cup_points species owner
## 0 0 7
## country_of_origin farm_name lot_number
## 1 359 1063
## mill ico_number company
## 315 151 209
## altitude region producer
## 226 59 231
## number_of_bags bag_weight in_country_partner
## 0 0 0
## harvest_year grading_date owner_1
## 47 0 7
## variety processing_method aroma
## 226 170 0
## flavor aftertaste acidity
## 0 0 0
## body balance uniformity
## 0 0 0
## clean_cup sweetness cupper_points
## 0 0 0
## moisture category_one_defects quakers
## 0 0 1
## color category_two_defects expiration
## 218 0 0
## certification_body certification_address certification_contact
## 0 0 0
## unit_of_measurement altitude_low_meters altitude_high_meters
## 0 230 230
## altitude_mean_meters
## 230
```

There are quite a few missing values and many columns have many. I will be just removing some of the columns with too many missing values, for instance `lot_number` and `farm_name`. Additionally, there will be removal of columns that do not heavily influence the goals of this project.

```
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.1.3

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.8      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.1
## v readr   2.1.2      v forcats 0.5.2

## Warning: package 'ggplot2' was built under R version 4.1.3

## Warning: package 'tibble' was built under R version 4.1.3

## Warning: package 'tidyr' was built under R version 4.1.3

## Warning: package 'readr' was built under R version 4.1.3

## Warning: package 'purrr' was built under R version 4.1.3

## Warning: package 'dplyr' was built under R version 4.1.3

## Warning: package 'stringr' was built under R version 4.1.3

## Warning: package 'forcats' was built under R version 4.1.3

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

Removal of columns

```
coffee = coffee_ratings%>%
  select(-farm_name, -lot_number, -mill, -ico_number, -altitude,
         -altitude_low_meters, -altitude_high_meters, -producer, -company,
         -expiration, -certification_address, -owner_1, -grading_date,
         -certification_contact, -unit_of_measurement)
```

```
#Remove Rows Containing Missing Values
```

```
coffee = na.omit(coffee)
```

```
#Changing Mass to all Imperial Units of Measurements
```

```
#selecting only items with lbs pattern within column to see how many
#Nathan F reminded me to the use of grep
coffee[grepl("lbs",coffee$bag_weight),]
```

```
## # A tibble: 18 x 28
##   total_~1 species owner count~2 region numbe~3 bag_w~4 in_co~5 harve~6 variety
##   <dbl> <chr>   <chr> <chr>   <chr>   <dbl> <chr>   <chr>   <chr>   <chr>
## 1    87.2 Arabica the ~ Costa ~ san r~    250 3 lbs Specia~ 2014 Caturra
## 2    86.3 Arabica fran~ Costa ~ west ~    250 2 lbs Specia~ 2015/2~ Caturra
## 3    85.3 Arabica the ~ Costa ~ west ~    250 3 lbs Specia~ 2014 Caturra
## 4    85.3 Arabica the ~ Costa ~ san r~    250 3 lbs Specia~ 2014 Caturra
## 5    84.7 Arabica fabi~ Costa ~ tarra~    50 1 lbs Specia~ 2014 Caturra
## 6    84.5 Arabica fabi~ Costa ~ tarra~    250 1 lbs Specia~ 2014 Caturra
## 7    83.8 Arabica germ~ United~ yauco~    18 5 lbs Specia~ 2013 Other
## 8    83.8 Arabica the ~ Guatem~ quetz~    250 3 lbs Specia~ 2012 Caturra
## 9    83.3 Arabica the ~ Costa ~ san r~    250 3 lbs Specia~ 2014 Caturra
## 10   83.3 Arabica itia~ Haiti thiot~     2 4 lbs Specia~ 2012 Typica
## 11    83 Arabica germ~ United~ yauco~    17 5 lbs Specia~ 2013 Other
## 12   81.5 Arabica myri~ Haiti dondo~    300 4 lbs Specia~ 2013 Blue M~
## 13   81.2 Arabica esse~ Guatem~ huehu~    36 55 lbs Bloss~ 2014 Pacama~
## 14   81.1 Arabica germ~ United~ yauco~    18 5 lbs Specia~ 2013 Other
## 15   80.9 Arabica chri~ Nicara~ matag~    275 1 lbs Specia~ 2013 Caturra
## 16   80.8 Arabica the ~ Costa ~ san r~    250 3 lbs Specia~ 2014 Caturra
## 17   79.3 Arabica the ~ Colomb~ pereir~    250 3 lbs Specia~ 2013 Caturra
## 18   79.1 Arabica germ~ United~ yauco~    18 5 lbs Specia~ 2013 Other
## # ... with 18 more variables: processing_method <chr>, aroma <dbl>,
## # flavor <dbl>, aftertaste <dbl>, acidity <dbl>, body <dbl>, balance <dbl>,
## # uniformity <dbl>, clean_cup <dbl>, sweetness <dbl>, cupper_points <dbl>,
## # moisture <dbl>, category_one_defects <dbl>, quakers <dbl>, color <chr>,
## # category_two_defects <dbl>, certification_body <chr>,
## # altitude_mean_meters <dbl>, and abbreviated variable names
## # 1: total_cup_points, 2: country_of_origin, 3: number_of_bags, ...
```

```
#separating out the columns based on the value and units associated with it
coffee = separate(data = coffee, col = bag_weight, into = c("weight", "type"), sep = " ")
```

```
#converted string to numeric
coffee$weight = as.numeric(coffee$weight)
```

```
#simple loop to change units
for(i in 1:length(coffee)){
  if(coffee[i,8]=="kg"){
    coffee[i,7] = round(coffee[i,7] * 2.20462,0)
    coffee[i,8] = "lbs"
  }
}
```

```
#remove type column as the weight col is uniform for unit type
coffee = coffee%>%
  select(-type)
```

```
#Changing Length to all Imperial Units of Measurements
```

```
#Note: If reshape lib is on, this will break
coffee = coffee%>%rename(avg_altitude=altitude_mean_meters)
coffee$avg_altitude = round(coffee$avg_altitude * 3.28084,0)
```

#Altering rows with years with form Year1/Year2 to the initial year (Year1)

```
coffee$harvest_year = substr(coffee$harvest_year,1,4)
coffee$harvest_year = as.numeric(coffee$harvest_year)
```

The above chunk was done do to the initial inception of that batch of coffee.

#Numerical Summary to see the data for potentail outliers

```
summary(coffee[,c(9,12:24,26,28)])
```

```
##  harvest_year      aroma      flavor      aftertaste      acidity
##  Min.   :2011    Min.   :5.080    Min.   :6.170    Min.   :6.170    Min.   :5.250
##  1st Qu.:2012    1st Qu.:7.420    1st Qu.:7.330    1st Qu.:7.170    1st Qu.:7.330
##  Median :2014    Median :7.580    Median :7.500    Median :7.420    Median :7.500
##  Mean   :2014    Mean   :7.559    Mean   :7.504    Mean   :7.374    Mean   :7.515
##  3rd Qu.:2015    3rd Qu.:7.750    3rd Qu.:7.670    3rd Qu.:7.580    3rd Qu.:7.670
##  Max.   :2018    Max.   :8.750    Max.   :8.670    Max.   :8.500    Max.   :8.580
##      body      balance      uniformity      clean_cup
##  Min.   :6.330    Min.   :6.080    Min.   : 6.000    Min.   : 0.000
##  1st Qu.:7.330    1st Qu.:7.330    1st Qu.:10.000    1st Qu.:10.000
##  Median :7.500    Median :7.500    Median :10.000    Median :10.000
##  Mean   :7.494    Mean   :7.488    Mean   : 9.871    Mean   : 9.849
##  3rd Qu.:7.670    3rd Qu.:7.670    3rd Qu.:10.000    3rd Qu.:10.000
##  Max.   :8.420    Max.   :8.580    Max.   :10.000    Max.   :10.000
##      sweetness      cupper_points      moisture      category_one_defects
##  Min.   : 1.33    Min.   :5.170    Min.   :0.00000    Min.   : 0.0000
##  1st Qu.:10.00    1st Qu.:7.250    1st Qu.:0.10000    1st Qu.: 0.0000
##  Median :10.00    Median :7.500    Median :0.11000    Median : 0.0000
##  Mean   : 9.93    Mean   :7.459    Mean   :0.09737    Mean   : 0.4262
##  3rd Qu.:10.00    3rd Qu.:7.670    3rd Qu.:0.12000    3rd Qu.: 0.0000
##  Max.   :10.00    Max.   :8.580    Max.   :0.17000    Max.   :31.0000
##      quakers      category_two_defects      avg_altitude
##  Min.   : 0.0000    Min.   : 0.000    Min.   :      3
##  1st Qu.: 0.0000    1st Qu.: 0.000    1st Qu.: 3609
##  Median : 0.0000    Median : 2.000    Median : 4300
##  Mean   : 0.1521    Mean   : 3.822    Mean   : 6145
##  3rd Qu.: 0.0000    3rd Qu.: 5.000    3rd Qu.: 5249
##  Max.   :11.0000    Max.   :47.000    Max.   :623898
```

The parameters for defects, quakers, and average altitude seem to have quite a range for values. Additionally, it can be seen for these fields that the max points are quite a ways away from the mean.


EDA / Visuals

```
library(ggplot2)
```


```
#Check for outliers in some of the fields
```

```
defect1_plt = ggplot(coffee, aes(y=category_one_defects)) +  
  geom_boxplot()  
defect2_plt = ggplot(coffee, aes(y=category_two_defects)) +  
  geom_boxplot()  
alt_plt = ggplot(coffee, aes(y=avg_altitude)) +  
  geom_boxplot()  
quakers = ggplot(coffee, aes(y=quakers)) +  
  geom_boxplot()
```


```
defect1_plt
```




```
defect2_plt
```



alt_plt



quakers



There are some outliers, but not that many that would result in a concern at this time. These fields may be removed from the current analysis due to the outliers and lack of variance within the data. As the majority of these values are 0. This will be removed in the upcoming data chunks. Additionally, as this project is to have more focus in analysis, there will be additional removal of fields. Specifically, the ownership items and their location details.

#Redfine the Dataset

```
c = coffee[,c(1:2,4,10:26,28)]
```


#Condense the data

```
c.v1 = c%>%pivot_longer(
  cols = !c(species, country_of_origin,variety,processing_method,color),
  names_to = "Variables",
  values_to = "Values")
```

Since, this data set will be re-used for other visuals. Otherwise the following code chunk could be used to generate a specific visual.

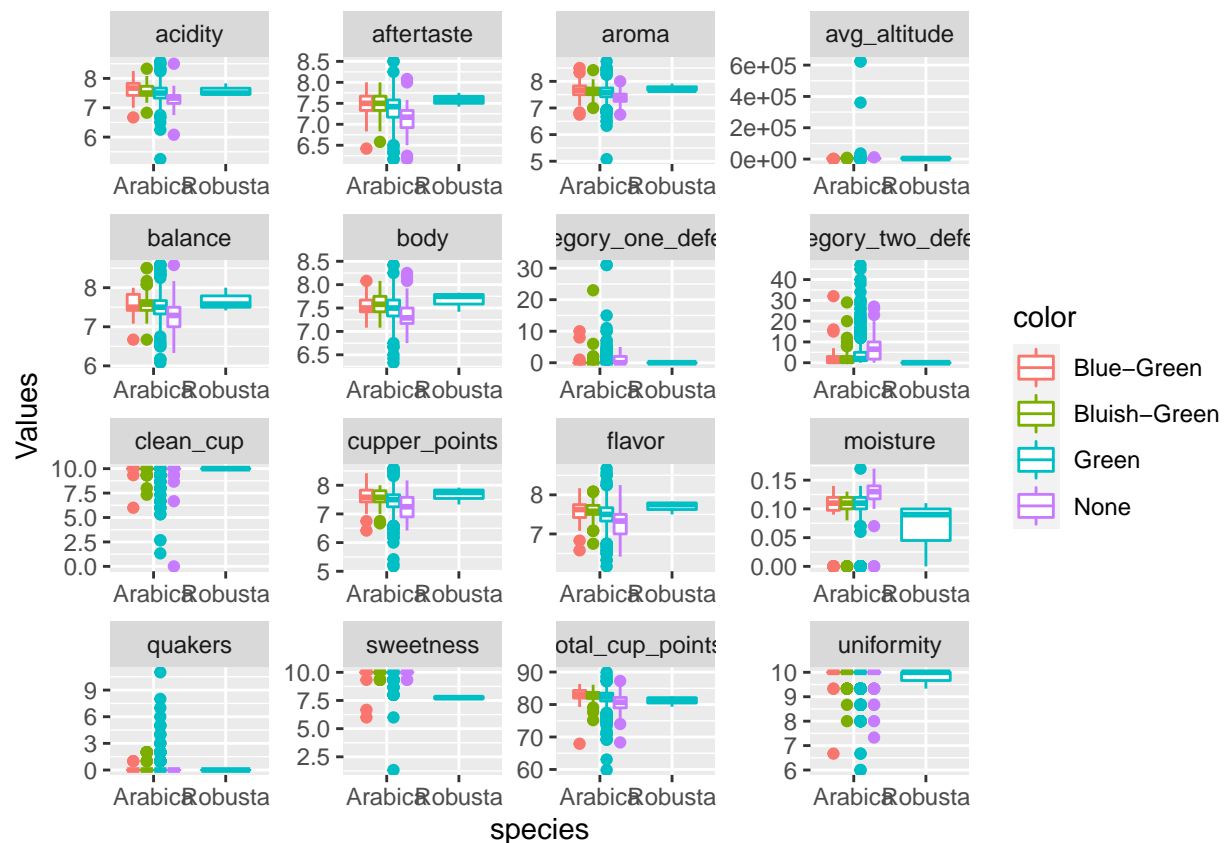
#Plot the data to see overall behavior

```
ggplot(c.v1,aes(x=species,y=Values))+geom_boxplot()+facet_wrap(~Variables,scales = "free")
```



#Plot the data to see overall behavior for specific field Coffee Color

```
ggplot(c.v1,aes(x=species,y=Values,color=color))+geom_boxplot()+facet_wrap(~Variables,scales = "free")
```

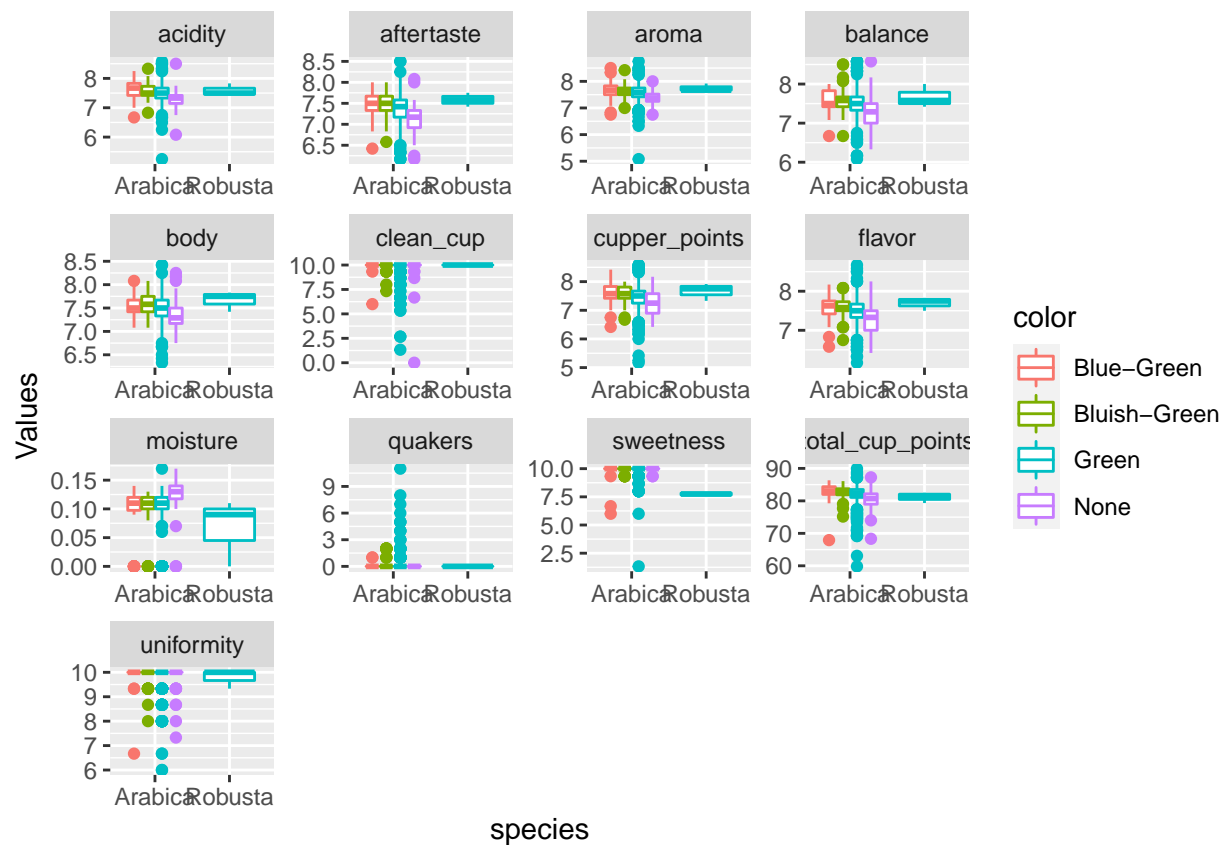


```
#Filter out the items that have known outliers
```

```
c.v2 = c.v1 %>%
  filter(Variables != 'avg_altitude' & Variables != 'category_one_defects' & Variables != 'category_two_defects')
```


```
#Re-run plot
```

```
ggplot(c.v2,aes(x=species,y=Values,color=color))+
  geom_boxplot()+
  facet_wrap(~Variables,scales = "free")
```




#How are the cup points distributed and where the 'weight' it is at by the Species and Coffee Color#

```
c.v2 %>%
  filter(Variables == 'total_cup_points')%>%
  ggplot(aes(x=species,y=Values,color=color))+
  geom_boxplot()
```

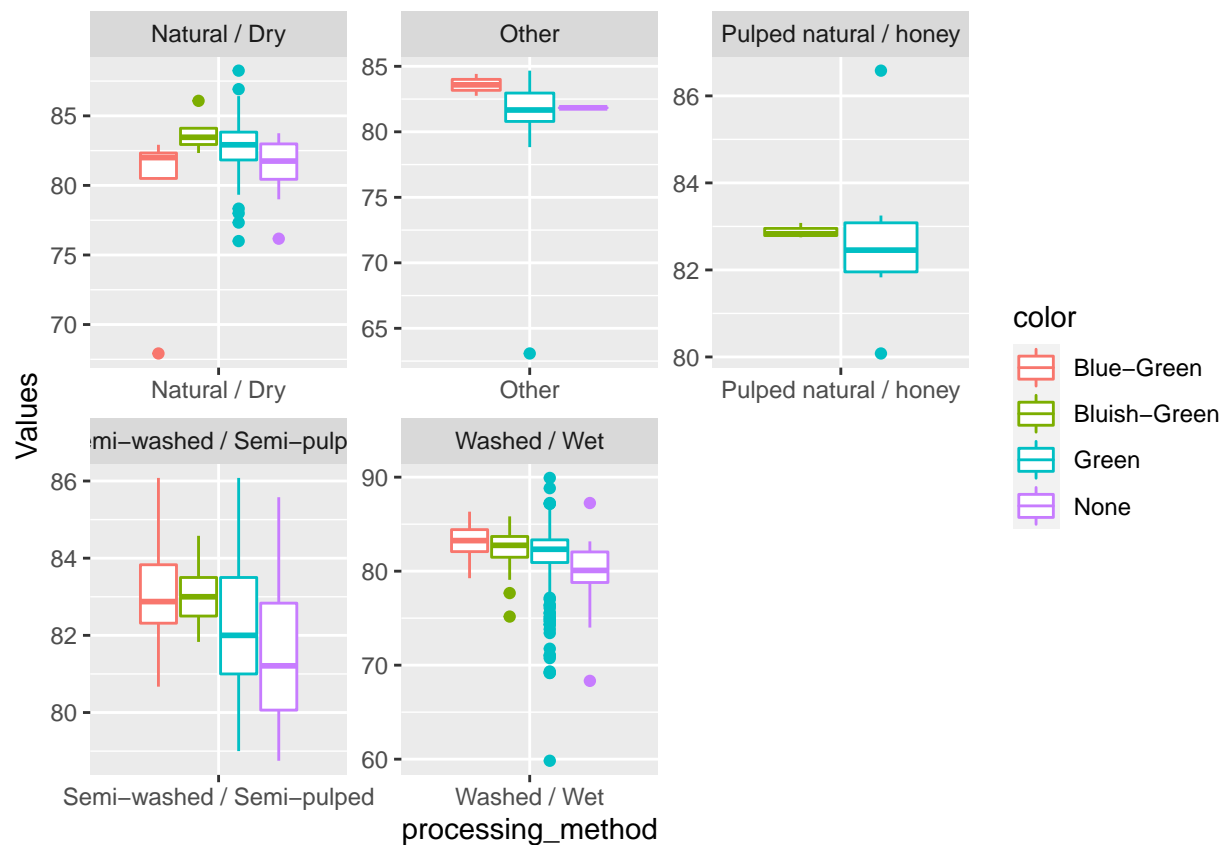


```
c.v2 %>%
  filter(Variables == 'total_cup_points')%>%
  ggplot(aes(x=species,y=Values,color=color))+
  geom_violin()
```

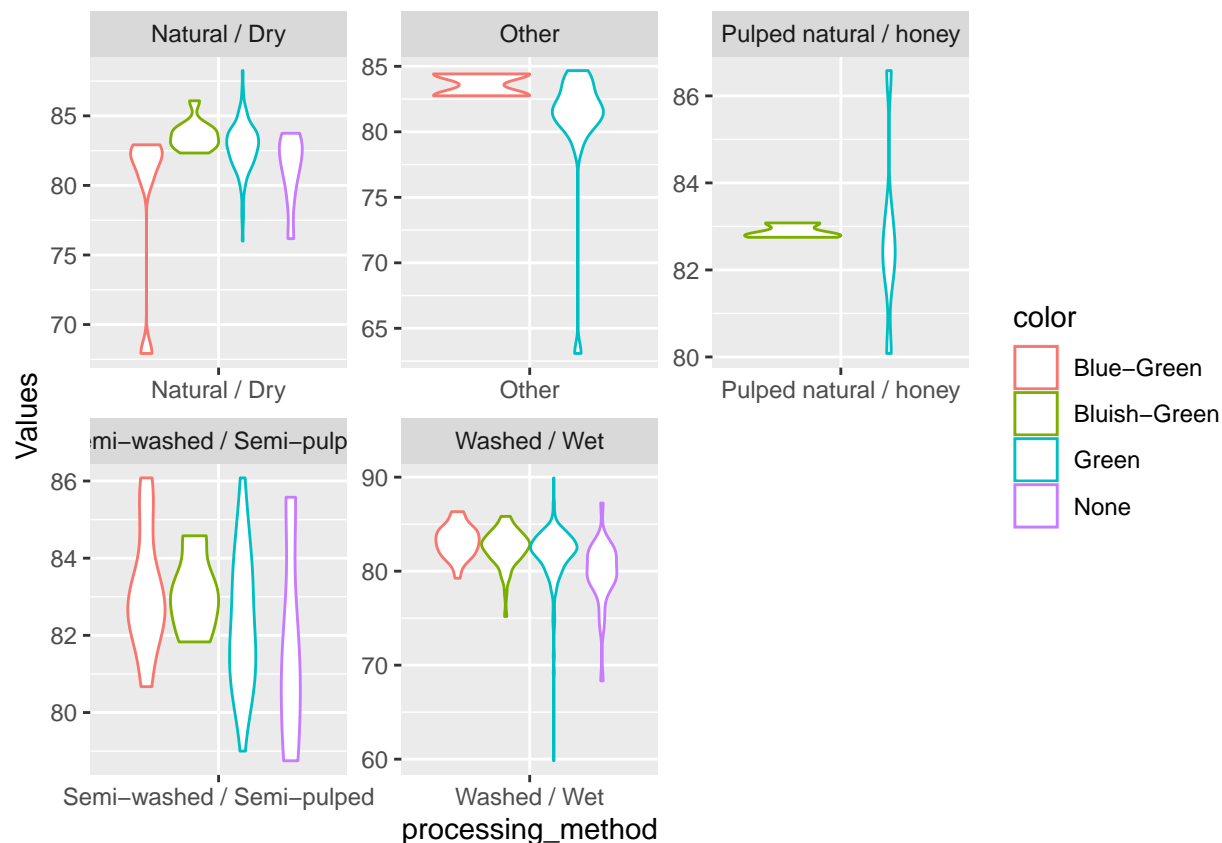


#How are the cup points distributed and where the 'weight' it is at by the Coffee Color and Processing Method

```
c.v2 %>%
  filter(Variables == 'total_cup_points')%>%
  ggplot(aes(x=processing_method,y=Values,color=color))+
  geom_boxplot()+
  facet_wrap(~processing_method,scales = "free")
```



```
c.v2 %>%
  filter(Variables == 'total_cup_points')%>%
  ggplot(aes(x=processing_method,y=Values,color=color))+
  geom_violin()+
  facet_wrap(~processing_method,scales = "free")
```





#Heatmap of Correlations

```
c = c[,c(1,6:16)]
cormat = cor(c)
melted = reshape::melt(cormat, varnames = c("ParameterX", "ParameterY"))
```

#Heatmap#

```
ggplot(data = melted, aes(x=ParameterX, y=ParameterY, fill=value)) +
  geom_tile(color = "white")+
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",
    midpoint = 0, limit = c(-1,1), space = "Lab",
    name="Pearson \n Correlation") +
  labs(x = "", y = "")+
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5), axis.text.x = element_text(angle = 45, vjust = 1, size = 12),
    coord_fixed())
```



#See the data make-up in a numerical summary

```
library(formattable)
```

Warning: package 'formattable' was built under R version 4.1.3

#Function for Calculating Frequency

```
freqq = function(df,col_i,col_j){
  a = df %>%
    group_by({col_i},{col_j}) %>%
    summarise(count = n()) %>%
    mutate(freq = formattable::percent(count / sum(count)))
  return(a)
}
```

#Overall Frequency all Countries

```
freqq(c.v1,Variables,Values)
```

'summarise()' has grouped output by 'Variables'. You can override using the
'.groups' argument.

A tibble: 611 x 4

```
## # Groups:   Variables [16]
##   Variables Values count freq
##   <chr>      <dbl> <int> <formttbl>
## 1 acidity    5.25     1 0.11%
## 2 acidity    6.08     1 0.11%
## 3 acidity    6.25     1 0.11%
## 4 acidity    6.5      1 0.11%
## 5 acidity    6.67     3 0.34%
## 6 acidity    6.75     2 0.22%
## 7 acidity    6.83     6 0.67%
## 8 acidity    6.92     7 0.78%
## 9 acidity     7      23 2.57%
## 10 acidity   7.08    25 2.80%
## # ... with 601 more rows
```

```
##Overall Frequency for Brazil
```

```
freq(c.v1%>%filter(country_of_origin=="Brazil"),Variables,Values)
```

```
## 'summarise()' has grouped output by 'Variables'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 216 x 4
## # Groups:   Variables [16]
##   Variables Values count freq
##   <chr>      <dbl> <int> <formttbl>
## 1 acidity    6.92     1 1.05%
## 2 acidity     7      1 1.05%
## 3 acidity    7.08     3 3.16%
## 4 acidity    7.17     4 4.21%
## 5 acidity    7.25     5 5.26%
## 6 acidity    7.33     8 8.42%
## 7 acidity    7.42     7 7.37%
## 8 acidity    7.5     26 27.37%
## 9 acidity    7.58     9 9.47%
## 10 acidity   7.67    13 13.68%
## # ... with 206 more rows
```

```
##Analysis Preparation
```

```
##Format new label (total_cup_points) to be categorical
```

```
coffee$tcp = coffee$total_cup_points
```

```
##Creating Bins for the Cup Points
```

```
for(i in 1:894){
  if(coffee[i,29] >= 80){
    coffee[i,29] = 80
  }
  else if(coffee[i,29] >= 70 & coffee[i,29] < 80){
    coffee[i,29] = 70
  }
}
```

```

}
else if(coffee[i,29] >= 60 & coffee[i,29] < 70){
  coffee[i,29] = 60
}
else{
  coffee[i,29] = 50
}
}
coffee$tcp = round(coffee$tcp,0)

```

While the bins could be more specific and look at every 2 or 5 points, it made more sense to use broader bins. This is due to trying to understand what makes a coffee from a specific bean have higher or lower overall cup points (i.e., what is the difference between 70s and 80s cup of coffee).

#Accuracy table for comparison between models

```

table_accuracy = matrix(nrow=4,ncol=1)
colnames(table_accuracy) = c('Accuracy')
rownames(table_accuracy) = c('DTree', 'NB', 'ANN', 'KNN')
table_accuracy

```

```

##      Accuracy
## DTree      NA
## NB        NA
## ANN       NA
## KNN       NA

```

This is to help determining which model or models is better than the others. If there are many with similar accuracy, then the model that is the easiest to interpret and explain to a general audience.

#Set seed so analysis is repeatable

```

set.seed(1)

```

For analysis

```

df = coffee[,c(9:22,25,29)]
for(i in 4 : 13){
  df[,i]=round(df[,i],2)
}

```

If the data was processing a bit slowly for initial predicting, as it was too granular so this step was helpful to making the ML run quicker.

#Fix issue with the Data#

```

df$processing_method= as.factor(df$processing_method)
df$variety = as.factor(df$variety)
df = df[,c(1:16)]
df$tcp = as.factor(df$tcp)
df$moisture = round(df$moisture,1)

```

This was missed earlier in the summary, but the fields that are characters, need to be changed to type factor for the analysis.

Simple k-fold cross validation(cv)

```
set.seed(1)
n = nrow(df)
folds = 10
tail = n%/%folds

rnd = runif(n)
rank = rank(rnd)

#block/chunk from cv
blk = (rank-1)%/%tail+1
blk = as.factor(blk)

#to see formation of folds
print(summary(blk))
```

```
##  1  2  3  4  5  6  7  8  9 10 11
## 89 89 89 89 89 89 89 89 89 89  4
```

Could turn the above into a more personalized cross validation method than one of the packages in an R library.

Predictive Analysis

#Decision Tree

```
library(rpart)
set.seed(1)

all.acc = numeric(0)
for(i in 1:folds){
  tree = rpart(tcp~.,df[blk != i,],method="class")
  pred = predict(tree,df[blk==i,],type="class")
  confMat = table(pred,df$tcp[blk==i])
  acc = (confMat[1,1]+confMat[2,2]+confMat[3,3]+confMat[4,4])/sum(confMat)
  all.acc = rbind(all.acc,acc)
}

print(mean(all.acc))
```

```
## [1] 0.9516854
```

```
table_accuracy[1,1] = mean(all.acc)
```

A 95% overall accuracy is really good! This indicates if following this tree, with details on a bean one could reasonable figure out what its overall score will be prior to evaluation. It also indicates what are the more important parameters are for a coffee scoring.

Example of a table matrix of predicted(rows) and actual(columns)

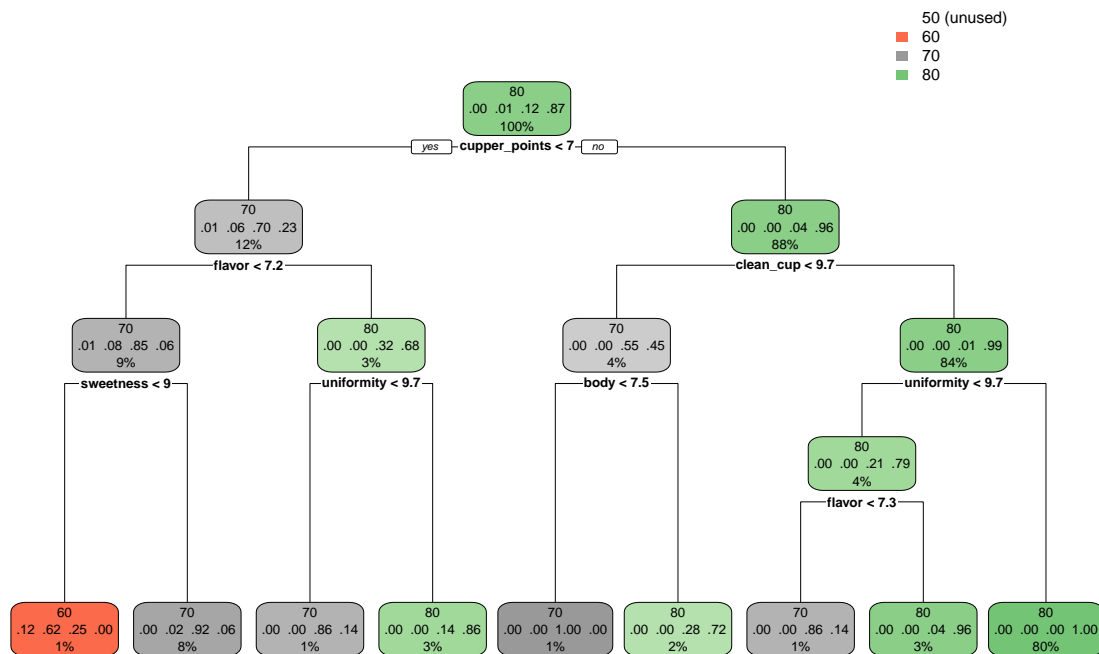
```
confMat
```

```
##
## pred 50 60 70 80
##    50  0  0  0  0
##    60  0  0  0  0
##    70  0  0 13  0
##    80  0  0  3 73
```

This indicates, for the given run, there were 3 miss classifications. Where the tree suggested that the bean should have been in the 80s, but was actually in the 70s.

Visual of Decision Tree

```
rpart.plot::rpart.plot(tree)
```



From this plot, I could just bin 50s with the 60sw group. This will help with future evaluations where re-binning the classifier would be a potential option to get more granular information.

Naive Bayes

```
library(e1071)

## Warning: package 'e1071' was built under R version 4.1.3

set.seed(1)

all.acc = numeric(0)
for(i in 1:folds){
  model = naiveBayes(tcp~.,df[blk != i,],method="class")
  pred = predict(model,df[blk==i,],type="class")
  confMat = table(pred,df$tcp[blk==i])
  acc = (confMat[1,1]+confMat[2,2]+confMat[3,3]+confMat[4,4])/sum(confMat)
  all.acc = rbind(all.acc,acc)
}

print(mean(all.acc))

## [1] 0.9550562

table_accuracy[2,1] = mean(all.acc)
```

Another nice and high accuracy for this PA!

Weird R Issue

```
#switch the classifier to numerical
df$tcp = round(as.numeric(df$tcp),0)
#them switch it back to a factor
df$tcp = as.factor(df$tcp)
```

This was a very weird issue. I knew that this was a factor was needed for the classifier. However, it was throwing a NaN for an accuracy value and just by switching the format back and forth corrected it.

Neural Network

```
library(nnet)
set.seed(1)

all.acc = numeric(0)
for(i in 1:folds){
  model = nnet(tcp~.,df[blk != i,], size = 11, trace=FALSE, rang=.06, decay=.006,maxit=500)
  pred = predict(model, df[blk==i,],type="class")
  confMat = table(factor(pred,levels=1:4),factor(df$tcp[blk==i],levels=1:4))
  acc = (confMat[1,1]+confMat[2,2]+confMat[3,3]+confMat[4,4])/sum(confMat)
  all.acc = rbind(all.acc,acc)
}
print(mean(all.acc))
```

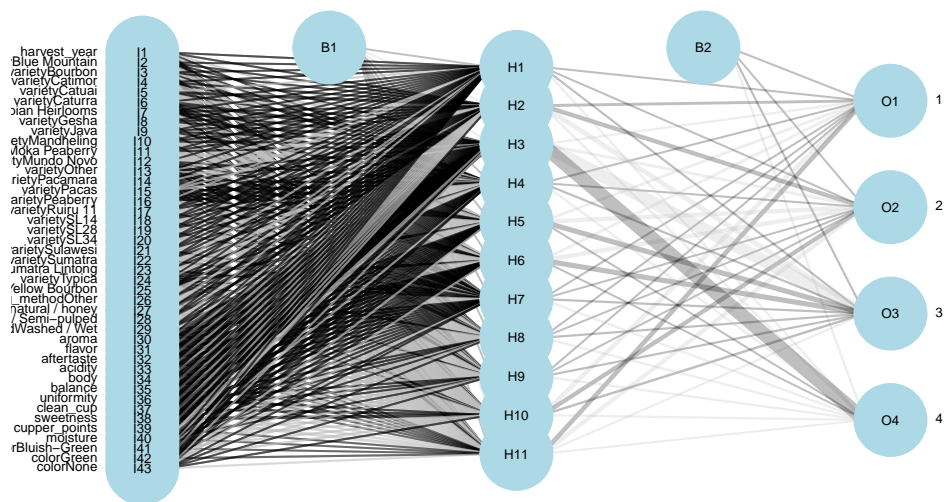
```
## [1] 0.8932584
```

```
table_accuracy[3,1] = mean(all.acc)
```

Not the best not the worst NN that I have seen. If there was more time, I would have liked to increase the classifiers and used a different library that allowed for more hidden layers.

Neural Network Visual

```
library("NeuralNetTools")  
plotnet(model, circle_cex=5, cex_val=.4, max_sp=TRUE, alpha_val=.25, skip=TRUE)
```



Note

An issue I ran in to:

I re-formatted the label/target field and went from a binary (good [>74]/bad [<75]) classifier to what is it currently; 50s, 60s, 70s, and 80s. However, when running the all of the PAs prior to neural network there were no strange issues. When running the NN I recieved an output accuracy of 0.003 an knew there was an issue.

There was an (un)interesting issue with NN table (well, all tables), as it was dropping the first two rows as it was not forward feeding into those nodes. The following is the work around to resolve this issue.

#Before

```
set.seed(1)
i=1
model = nnet(tcp~.,df[blk != i,], size = 10, trace=FALSE, wgt=.05)
pred = predict(model, df[blk==i,],type="class")
confMat = table(pred,df$tcp[blk==i])
confMat
```

```
##
## pred  1  2  3  4
##      3  1  0 16 72
```

#After

```
set.seed(1)
i=1
model = nnet(tcp~.,df[blk != i,], size = 10, trace=FALSE, wgt=.05)
pred = predict(model, df[blk==i,],type="class")
confMat = table(factor(pred,levels=1:4),factor(df$tcp[blk==i],levels=1:4))
confMat
```

```
##
##      1  2  3  4
##  1  0  0  0  0
##  2  0  0  0  0
##  3  1  0 16 72
##  4  0  0  0  0
```

This was then applied to all of the PAs.

K-Nearest Neighbor Preparation

```
set.seed(1)
df$tcp = as.factor(df$tcp)
library (caret)
```

```
## Warning: package 'caret' was built under R version 4.1.3
```

```
## Loading required package: lattice
```

```
##
## Attaching package: 'caret'
```


```
## The following object is masked from 'package:purrr':
##
## lift
```

```
trControl <- trainControl(method = "cv", number = 10)
knn = df[,]
```

KNN

```
set.seed(1)
model <- train(tcp ~ .,
               method = "knn",
               tuneGrid = expand.grid(k = 1:10),
               trControl = trControl,
               data = knn)
acc = mean(model$results$Accuracy)
table_accuracy[4,1] = acc

plot(model)
```



This is a visual to see how many neighbors the KNN will be running. From this visual it could possibly run at 9 groups due to the accuracy level.

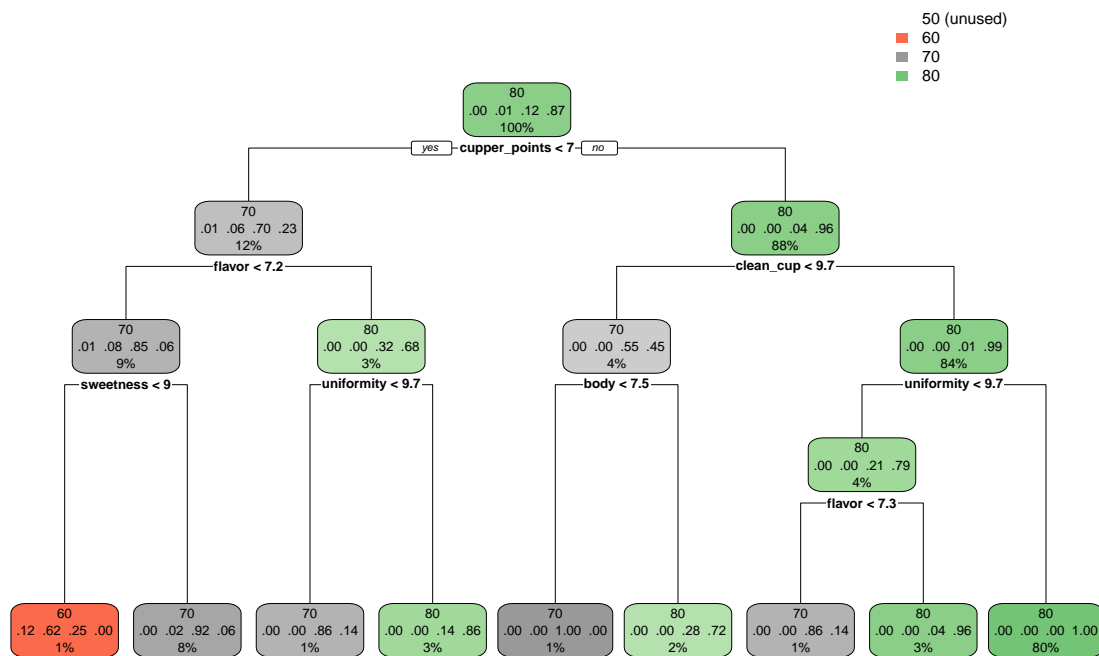
#View Accuracy Table

```
tab = round(table_accuracy,4)
tab
```

```
##      Accuracy
## DTree  0.9517
## NB     0.9551
## ANN    0.8933
## KNN    0.9325
```

```
##Preferred Model
```

```
rpart.plot::rpart.plot(tree)
```



Top 3 parameters for understanding a coffee's score.

~Cupper points are the most informative parameter in deciding if a coffee is to be in the 80s or below this.

~If place coffee is <7 cupper points, the next deciding factor is how good is the flavor of the coffee.

~ If coffee is >7 cupper points, the next deciding factor is how clean the coffee leaves the cup.

For further analysis

```
df2 = coffee[,c(4,5,9:22,25,29)]
for(i in 6 : 16){
  df2[,i]=round(df2[,i],2)

df2$processing_method= as.factor(df2$processing_method)
df2$variety = as.factor(df2$variety)
df2$tcp = as.factor(df2$tcp)
df2$moisture = round(df2$moisture,1)
```

```
df2$color = as.factor(df2$color)
df2$country_of_origin = as.factor(df2$country_of_origin)
df2$region = as.factor(df2$region)
df3 = df2[,c(1,3:18)]
}
```

```
set.seed(1)
n = nrow(df3)
folds = 10
tail = n%%folds

rnd = runif(n)
rank = rank(rnd)

#block/chunk from cv
blk = (rank-1)%/%tail+1
blk = as.factor(blk)

#to see formation of folds
print(summary(blk))
```

```
##  1  2  3  4  5  6  7  8  9 10 11
## 89 89 89 89 89 89 89 89 89 89  4
```

```
set.seed(1)

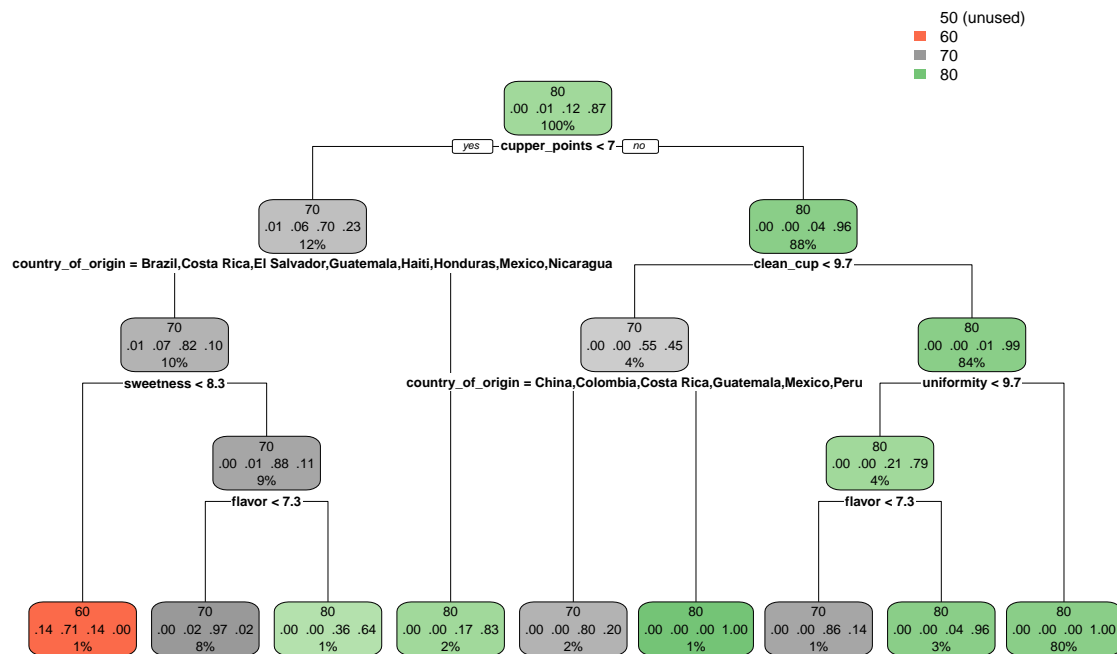
all.acc = numeric(0)
for(i in 1:folds){
  tree = rpart(tcp~.,df3[blk != i,],method="class")
  pred = predict(tree,df3[blk==i,],type="class")
  confMat = table(pred,df3$tcp[blk==i])
  acc = (confMat[1,1]+confMat[2,2]+confMat[3,3]+confMat[4,4])/sum(confMat)
  all.acc = rbind(all.acc,acc)
}

print(mean(all.acc))
```

```
## [1] 0.947191
```

Interestingly, adding countries lowers the accuracy.

```
rpart.plot::rpart.plot(tree)
```



From the visual, it appears that Central and South America do not produce good coffee.